



Women in Parliament – Tidyverse Edition

Saghir Bashir

This version was compiled on April 7, 2021

We will use the World Bank's indicator data for "Women in Parliament" as a case study when working with the tidyverse suite of R packages. We will guide you through the geographical and time trends for the percentage of women in national parliaments. We will start by learning about and understanding the raw data, which we will then process ("wrangle") in preparation for some exploratory analysis.

Women in Parliament | World Bank Indicator | Tidyverse | dplyr | tidyr | ggplot2

1. Preface

We present a real-life case study for some of the tidyverse¹ package using the World Bank's "Women in Parliament" indicator data. To get the most out of this case-study guide, repeat the examples and do the exercises whilst reading it.

Guide materials. You can download materials for this guide from this link:

- <https://ilustat.com/shared/WiP-tidyverse.zip>

Unzip the file, which contains the data, this guide and an R script exercise file. We advise you to work with "WiP-Exercise.R" file to follow the examples and do the exercises. If you are using RStudio, you can double click on "WiP-tv.Rproj" to get started.

2. Objectives

Explore the geographical and time trends for the percentage of women² in national parliaments.

3. Understanding the Data

The World Bank Data. The raw data for "Proportion of seats held by women in national parliaments" includes the percentage of women in parliament ("single or lower parliamentary chambers only") by country (region) and year. It can be downloaded from:³

- <https://data.worldbank.org/indicator/SG.GEN.PARL.ZS>

As part of its "open data" mission the World Bank offers "free and open access to global development data" kindly licensed under the "Creative Commons Attribution 4.0 (CC-BY 4.0)".⁴

Source Data. The data originates from the "Inter-Parliamentary Union" (IPU)⁵ which provides an "Archive of statistical data on the percentage of women in national parliaments" going back to 1997 on a monthly basis:

- <http://archive.ipu.org/wmn-e/classif-arc.htm>

The World Bank data is for "single or lower parliamentary chambers only", while the IPU also presents data for "Upper Houses or Senates". Moreover, the IPU provides the actual numbers used to calculate the percentages (which the World Bank does not).

Data limitations. Take caution when interpreting these data, as parliamentary systems vary from country to country, and in some cases over time. Some of the issues to consider include:

- Who has, and who does not have, the right to become a Member of Parliament (MP)?
- How does someone become an MP? Through democratic elections? How is "democratic election" defined?
- What is the real power of MPs and their parliament? Can MPs make a difference?

Data definitions & assumptions.

"Women". The definition for "women" is not given, so we will assume that it refers to a binary classification for gender (sex).

"Country (Region)". The definition of countries and regions can change over time. (e.g. formation of new countries after conflicts, new member states joining a pre-existing collective). How are these changes reflected in the data? How do they affect the interpretation?

Pro tip. Understand the limitations of your data before anybody else points them out to you.

4. About the data file

The data is stored in a file called:

- `API_SG.GEN.PARL.ZS_DS2_en_csv_v2_2163427.csv`

To simplify things we have copied it to `WB-WiP.csv` (which also allows us to maintain the original file in case something goes wrong).

Pro tip. Always keep a backup copy of the data. Alternatively, set the data file(s) to "read-only" to protect it from being overwritten or modified.

Exercise. It is important to look at and understand the contents of the file before you start using it. Using a text editor or a spreadsheet software, open the `WB-WiP.csv` file (in the data directory). What do you observe in the contents of this file?

Content and Structure. The first four lines of `WB-WiP.csv` can be ignored, since they contain two lines of meta-information and two blank lines, as follows:

```
1 "Data Source","World Development Indicators",
2
3 "Last Updated Date","2021-03-19",
4
```

The fifth line contains the column (variable) names and the body of data starts in the sixth line. It is important to note that there was no collection of data for a majority of the years, which means that it is "missing".

5. Loading tidyverse packages

We will load the tidyverse packages we plan to use individually (messages have been suppressed).⁶

⁶We could have used `library(tidyverse)` but we prefer to load packages individually and only those that we will use.

¹For more information on the tidyverse see <https://www.tidyverse.org/>.

²The objective could be termed neutrally as "gender trends" but we will keep it per the World Bank data.

³The `wbstats` R package (<https://cran.r-project.org/web/packages/wbstats/>) gives access to a "tidier" version of the World Bank indicator data.

⁴<https://datacatalog.worldbank.org/public-licenses/cc-by>.

⁵Inter-Parliamentary Union: <https://www.ipu.org/>.

```
library(here)
library(readr)
library(dplyr)
library(tidyr)
library(ggplot2)
library(gghighlight)
```

6. Importing the data

Based on our findings above, we can “skip” the first four lines and treat the fifth line as column (variable) names.

```
wip <- read_csv(here("data", "WB-WiP.csv"),
               skip = 4)
```

```
# Warning: Missing column names filled in:
# 'X66' [66]
```

Messages. We have suppressed some of the messages but left the one about variable X66 (we will come back to it below).

Exercise. Check what you have read by typing “wip” in the console window. What do you observe? Type “class(wip)” and “glimpse(wip)” to confirm that “wip” is of class “tbl_df”.

“Fix” column names. Some of the column names contain spaces while others are numeric:

```
head(names(wip))
# [1] "Country Name" "Country Code"
# [3] "Indicator Name" "Indicator Code"
# [5] "1960" "1961"
tail(names(wip))
# [1] "2016" "2017" "2018" "2019" "2020" "X66"
```

By using the `make.names()` function we don’t need to use back ticks (‘) around the column names (e.g. ‘col name’).

```
names(wip) <- make.names(names(wip))
head(names(wip))
# [1] "Country.Name" "Country.Code"
# [3] "Indicator.Name" "Indicator.Code"
# [5] "X1960" "X1961"
tail(names(wip))
# [1] "X2016" "X2017" "X2018" "X2019" "X2020"
# [6] "X66"
```

7. Data Wrangling Aims

We can simplify the production of summaries and plots by restructuring the current wip dataset (which has 66 columns) to the following format:

Country	Year	pctWiP
Country AAA	1997	##.##
Country AAA	1998	##.##
Country AAA	1999	##.##
...		

pctWiP refers to the percentage of women in parliament.

Key information retained. These three columns will contain the same information as the wip dataset but in a more usable format. We will also add a variable for the ratio of male to female MPs.

Superfluous columns. We will start by removing columns X66, Indicator.Name and Indicator.Code. There are years without any data but they will be removed automatically later (when restructuring from “wide” to “long” format).

Column X66 is created automatically due to an extra comma at the end of the column names (fifth) line of WB-WiP.csv:

```
... , "2017", "2018", "2019", "2020",
```

Check. Before removing it check that all values are NA.

```
wip %>% pull(X66) %>% is.na(.) %>% all(.)
# [1] TRUE
```

Column Indicator.Name has the unique value “Proportion of seats held by women in national parliaments (%)” and in Indicator.Code it is “SG.GEN.PARL.ZS”. As there is only one indicator in this dataset we will remove these two columns.

Exercise. Confirm that both Indicator.Name and Indicator.Code have the same values for all observations. Hint: Use either `count()` or `distinct()` functions.

Removing columns. The indicator and X66 columns can be removed. We will also rename “Country.Name” as “Country” and “Country.Code” as “Code”.

```
wip2 <- wip %>%
  select(-Indicator.Name, -Indicator.Code,
        -X66) %>%
  rename(Country=Country.Name, Code=Country.Code)
head(names(wip2))
# [1] "Country" "Code" "X1960" "X1961"
# [5] "X1962" "X1963"
tail(names(wip2))
# [1] "X2015" "X2016" "X2017" "X2018" "X2019"
# [6] "X2020"
```

Reshape to long format. We want to transform the data so that for each country the year (column) data becomes a row. At the same time we will remove the missing data (with the `na.rm` option). We will also create a numeric Year variable and a Ratio of men to women in parliament.

```
WP <- wip2 %>%
  pivot_longer(starts_with("X"),
               names_to = "YearC",
               values_to = "pctWiP",
               values_drop_na = TRUE) %>%
  mutate(Year = parse_number(YearC),
         Ratio = (100-pctWiP)/pctWiP) %>%
  select(Country, Code, Year, pctWiP, Ratio) %>%
  arrange(Country, Year)
# Look at the contents of WP
glimpse(WP)
# Rows: 5,391
# Columns: 5
# $ Country <chr> "Afghanistan", "Afghanistan~
# $ Code <chr> "AFG", "AFG", "AFG", "AFG", ~
# $ Year <dbl> 2005, 2006, 2007, 2008, 200~
# $ pctWiP <dbl> 27.3, 27.3, 27.7, 27.7, 27.~
# $ Ratio <dbl> 2.66, 2.66, 2.61, 2.61, 2.6~
```

8. Questions

The objective is to look at the geographical and time trends in the data. We will answer the following questions.

- What are the time trends for Portugal?
- How does Portugal compare to other countries?
- Which countries have the highest percentage of women in parliament by year?
- How do continents compare?
- What are the global trends over time?

Exercise - Without Programming.

- Which country do you think has the highest percentage of women in parliament?
- In each continent (i.e. Africa, Americas, Asia, Europe and Oceania), which country has the highest percentage of women in parliament?
- What is the world percentage of women in parliament in 2020?

9. Exploratory Analysis


Select a country. This guide explores how Portugal performs over time and compared to other countries. Before continuing, select another country for yourself to repeat the examples and do the exercises.

Time trends for Portugal. First look at the raw data.

```
# Reset tibble print option to see more rows
options(tibble.print_max = 25)
WP %>% filter(Country=="Portugal")
# # A tibble: 24 x 5
#   Country Code Year pctWiP Ratio
#   <chr>   <chr> <dbl> <dbl> <dbl>
# 1 Portugal PRT 1997 13.0 6.67
# 2 Portugal PRT 1998 13.0 6.67
# 3 Portugal PRT 1999 18.7 4.35
# 4 Portugal PRT 2000 17.4 4.75
# 5 Portugal PRT 2001 18.7 4.35
# 6 Portugal PRT 2002 19.1 4.23
# 7 Portugal PRT 2003 19.1 4.23
# 8 Portugal PRT 2004 19.1 4.23
# 9 Portugal PRT 2005 21.3 3.69
# 10 Portugal PRT 2006 21.3 3.69
# 11 Portugal PRT 2007 28.3 2.54
# 12 Portugal PRT 2008 28.3 2.54
# 13 Portugal PRT 2009 27.4 2.65
# 14 Portugal PRT 2010 27.4 2.65
# 15 Portugal PRT 2011 28.7 2.48
# 16 Portugal PRT 2012 28.7 2.48
# 17 Portugal PRT 2013 31.3 2.19
# 18 Portugal PRT 2014 31.3 2.19
# 19 Portugal PRT 2015 34.8 1.87
# 20 Portugal PRT 2016 34.8 1.87
# 21 Portugal PRT 2017 34.8 1.87
# 22 Portugal PRT 2018 34.8 1.87
# 23 Portugal PRT 2019 38.7 1.58
# 24 Portugal PRT 2020 40 1.5
```

Visualisation. It is easier to find trends within a plot.

```
WP %>%
  filter(Country=="Portugal") %>%
  ggplot(aes(Year, pctWiP)) +
  geom_line() + geom_point() +
  scale_y_continuous(limits=c(0, 50)) +
  ylab("% Women in Parliament")
```




Interpretation. In 1997 Portugal had 13% women in parliament (i.e. 6.67 men for each woman), which increased to 40% (i.e. 1.5 men for each woman) in 2020. This still falls short of 50% (i.e. point of gender parity in parliament).

Exercise. For your chosen country look at the time trend data and the plot. What is your interpretation? How does it compare to Portugal?

Portugal versus European Union (EU) countries. We selected six EU countries (due to space limitations) for comparison. It would be better to compare all EU and/or all European countries.

```
WP %>%
  filter(Country %in% c("Portugal", "Sweden",
    "Spain", "Hungary", "Romania", "Finland",
    "Germany", "European Union")) %>%
  ggplot(aes(Year, pctWiP, colour=Country)) +
  geom_line() +
  geom_point() +
  scale_x_continuous(breaks=seq(1995, 2020, 5)) +
  scale_y_continuous(limits=c(0, 50),
    breaks=seq(0, 50, by=10)) +
  ggtitle("Women in Parliament: EU Countries") +
  ylab("% Women in Parliament")
```



Interpretation. Since 2007 Portugal has had more women in parliament than the European Union average. The key point to note is that none of these countries reaches equality between males and females in parliament, although Sweden and Finland come closest.

A couple of points to note.

“Germany”. In October 1997, the process of “German reunification” led to the creation of Germany, which united the former “German Democratic Republic” (East Germany) and the “Federal Republic of Germany” (West Germany). Therefore, since reunification, the data is presented for the reunified “Germany” only. Careful thought should be given to handling, analysing and interpreting any pre-reunification data (if available).

“European Union”. The “European Union” has changed over time (unlike the “continent of Europe”). It started in the 1950s as a block of six European countries (known as the “European Community”) and has expanded over the years to 28 countries (with the United Kingdom having departed now). This raises the question of how the European Union average is calculated. For a given year, is it calculated based on the actual member states in that year or on all of the current member states?

Exercises. Compare the country of your choice to four or five other countries by plotting a line graph similar to the one above.

Countries with the highest percentage of women in parliament.

A quick answer can be obtained by looking at the highest percentages.

```
WP %>%
  arrange(-pctWiP) %>%
  head(10)
# # A tibble: 10 x 5
#   Country Code Year pctWiP Ratio
#   <chr> <chr> <dbl> <dbl> <dbl>
# 1 Rwanda RWA 2013 63.8 0.569
# 2 Rwanda RWA 2014 63.8 0.569
# 3 Rwanda RWA 2015 63.8 0.569
# 4 Rwanda RWA 2016 63.8 0.569
# 5 Rwanda RWA 2017 61.2 0.633
# 6 Rwanda RWA 2018 61.2 0.633
# 7 Rwanda RWA 2019 61.2 0.633
# 8 Rwanda RWA 2020 61.2 0.633
# 9 Rwanda RWA 2008 56.2 0.778
# 10 Rwanda RWA 2009 56.2 0.778
```

Data speaks. Are you surprised? Data can be very enlightening.

Highest percentage by year. Which countries have the highest percentage of women in parliament by year?

```
WP %>%
  group_by(Year) %>%
  arrange(Year, -pctWiP) %>%
  filter(row_number()==1)
# # A tibble: 24 x 5
#   # Groups:   Year [24]
#   Country Code Year pctWiP Ratio
#   <chr> <chr> <dbl> <dbl> <dbl>
# 1 Sweden SWE 1997 40.4 1.48
# 2 Sweden SWE 1998 40.4 1.48
# 3 Sweden SWE 1999 42.7 1.34
# 4 Sweden SWE 2000 42.7 1.34
# 5 Sweden SWE 2001 42.7 1.34
# 6 Sweden SWE 2002 45.0 1.22
# 7 Rwanda RWA 2003 48.8 1.05
# 8 Rwanda RWA 2004 48.8 1.05
# 9 Rwanda RWA 2005 48.8 1.05
# 10 Rwanda RWA 2006 48.8 1.05
# 11 Rwanda RWA 2007 48.8 1.05
```

```
# 12 Rwanda RWA 2008 56.2 0.778
# 13 Rwanda RWA 2009 56.2 0.778
# 14 Rwanda RWA 2010 56.2 0.778
# 15 Rwanda RWA 2011 56.2 0.778
# 16 Rwanda RWA 2012 56.2 0.778
# 17 Rwanda RWA 2013 63.8 0.569
# 18 Rwanda RWA 2014 63.8 0.569
# 19 Rwanda RWA 2015 63.8 0.569
# 20 Rwanda RWA 2016 63.8 0.569
# 21 Rwanda RWA 2017 61.2 0.633
# 22 Rwanda RWA 2018 61.2 0.633
# 23 Rwanda RWA 2019 61.2 0.633
# 24 Rwanda RWA 2020 61.2 0.633
```

Merging continent. The variable Country in the WP dataset is a mix of countries and regions (e.g. “European Union”, “South Asia” and “World”). To present the highest percentages grouped by continent we need to add it. Luckily, given the large number of R packages available, we can merge the “continent” from the “codelist” dataset in the “countrycode” package.

```
# Ensure that 'countrycode' package is installed.
# install.packages("countrycode")
library(countrycode)
cl <- codelist %>%
  select(continent, wb) %>%
  rename(Code = wb, Continent = continent)
cWP <- WP %>%
  left_join(cl, by = "Code")
```

Highest percentages by year and continent. Which countries have the highest percentages in 1997 and 2020?

```
cWP %>%
  filter(Year %in% c(1997, 2020) &
         !is.na(Continent)) %>%
  group_by(Continent, Year) %>%
  arrange(Continent, Year, -pctWiP) %>%
  filter(row_number()==1) %>%
  select(Continent, Country, Year, pctWiP, Ratio)
# # A tibble: 10 x 5
#   # Groups:   Continent, Year [10]
#   Continent Country Year pctWiP Ratio
#   <chr> <chr> <dbl> <dbl> <dbl>
# 1 Africa Seychelles 1997 27.3 2.67
# 2 Africa Rwanda 2020 61.2 0.633
# 3 Americas Argentina 1997 27.6 2.62
# 4 Americas Cuba 2020 53.2 0.879
# 5 Asia Vietnam 1997 26.2 2.81
# 6 Asia United Arab ~ 2020 50 1
# 7 Europe Sweden 1997 40.4 1.48
# 8 Europe Sweden 2020 47.0 1.13
# 9 Oceania New Zealand 1997 29.2 2.43
# 10 Oceania New Zealand 2020 40.8 1.45
```

Decline in percentage. Which countries have had a decline in percentage since their first measurement (not always 1997)?

```
dWP <- cWP %>%
  group_by(Country) %>%
  arrange(Country, Year) %>%
  filter(row_number()==1 | row_number()==n()) %>%
  mutate(pctDiff = pctWiP -
         lag(pctWiP, order_by=Country)) %>%
  filter(pctDiff<0 & !is.na(Continent)) %>%
```



```

arrange(pctDiff)
dWP %>% select(Country, pctDiff)
# # A tibble: 11 x 2
# # Groups:   Country [11]
#   Country          pctDiff
#   <chr>            <dbl>
# 1 Seychelles        -6.06
# 2 Korea, Dem. Peoples Rep. -2.47
# 3 Bahamas, The     -2.18
# 4 Tuvalu            -2.08
# 5 Papua New Guinea  -1.83
# 6 Maldives          -1.65
# 7 Haiti             -1.07
# 8 Congo, Rep.       -0.742
# 9 Yemen, Rep.       -0.332
# 10 Afghanistan      -0.293
# 11 Oman              -0.0841

```

Ranked status. Another way to look at the data is to look at the ranking of countries, which could be done at a global level or by continent. Nonetheless, the results should be interpreted with caution and an understanding of the actual percentages. For example, if most countries were around the 50% mark, rankings could be misleading and subject to random fluctuations.

Global ranks by year. We will rank the countries by year based on the percentage of women in parliaments. The countries with the highest percentage will be ranked first and the lowest last. A total for the number of countries with data is included as it varies by year.

```

cWPrankG <- cWP %>%
  filter(!is.na(Continent)) %>%
  group_by(Year) %>%
  mutate(RankG = rank(-pctWiP),
         TotalG = n())

```

Global ranking – Portugal.

```

cWPrankG %>%
  filter(Country=="Portugal") %>%
  select(Country, Year, pctWiP, Ratio, RankG,
         TotalG) %>%
  arrange(Year)
# # A tibble: 24 x 6
# # Groups:   Year [24]
#   Country Year pctWiP Ratio RankG TotalG
#   <chr>   <dbl> <dbl> <dbl> <dbl> <int>
# 1 Portugal 1997 13.0 6.67 41.5 160
# 2 Portugal 1998 13.0 6.67 49.5 162
# 3 Portugal 1999 18.7 4.35 24 153
# 4 Portugal 2000 17.4 4.75 34 157
# 5 Portugal 2001 18.7 4.35 33 167
# 6 Portugal 2002 19.1 4.23 42 161
# 7 Portugal 2003 19.1 4.23 46 174
# 8 Portugal 2004 19.1 4.23 54 181
# 9 Portugal 2005 21.3 3.69 46 185
# 10 Portugal 2006 21.3 3.69 50 189
# 11 Portugal 2007 28.3 2.54 28 188
# 12 Portugal 2008 28.3 2.54 28 187
# 13 Portugal 2009 27.4 2.65 33 187
# 14 Portugal 2010 27.4 2.65 34 187
# 15 Portugal 2011 28.7 2.48 31 188
# 16 Portugal 2012 28.7 2.48 35 189
# 17 Portugal 2013 31.3 2.19 36 184
# 18 Portugal 2014 31.3 2.19 36 187

```

```

# 19 Portugal 2015 34.8 1.87 29 189
# 20 Portugal 2016 34.8 1.87 27 192
# 21 Portugal 2017 34.8 1.87 28 190
# 22 Portugal 2018 34.8 1.87 29 191
# 23 Portugal 2019 38.7 1.58 28 192
# 24 Portugal 2020 40 1.5 22.5 190

```

Interpretation. Portugal has generally been ranked in the first quartile (25%) of countries in the world, with the fluctuations of its ranking most likely due to random variation.

Exercise. For your chosen country, interpret its ranking over the years. How does it compare to Portugal?

Continent ranks by year. We will rank the countries by year within a continent based on the percentage of women in parliaments. The countries with the highest percentage will be ranked first and the lowest last. A total for the number of countries with data, within each continent, is included as it varies by year.

```

cWPx <- cWPrankG %>%
  filter(!is.na(Continent)) %>%
  group_by(Continent, Year) %>%
  mutate(RankC = rank(-pctWiP),
         TotalC = n())

```

Portugal's ranking in Europe.


```

cWPx %>%
  ungroup() %>%
  filter(Country=="Portugal") %>%
  select(Country, Year, pctWiP, Ratio, RankC,
         TotalC) %>%
  arrange(Year)
# # A tibble: 24 x 6
#   Country Year pctWiP Ratio RankC TotalC
#   <chr>   <dbl> <dbl> <dbl> <dbl> <int>
# 1 Portugal 1997 13.0 6.67 16.5 38
# 2 Portugal 1998 13.0 6.67 18.5 36
# 3 Portugal 1999 18.7 4.35 12 37
# 4 Portugal 2000 17.4 4.75 16 37
# 5 Portugal 2001 18.7 4.35 15 40
# 6 Portugal 2002 19.1 4.23 15 38
# 7 Portugal 2003 19.1 4.23 17 40
# 8 Portugal 2004 19.1 4.23 21 41
# 9 Portugal 2005 21.3 3.69 20 41
# 10 Portugal 2006 21.3 3.69 21 43
# 11 Portugal 2007 28.3 2.54 15 43
# 12 Portugal 2008 28.3 2.54 13 43
# 13 Portugal 2009 27.4 2.65 15 43
# 14 Portugal 2010 27.4 2.65 15 43
# 15 Portugal 2011 28.7 2.48 14 43
# 16 Portugal 2012 28.7 2.48 15 43
# 17 Portugal 2013 31.3 2.19 16 43
# 18 Portugal 2014 31.3 2.19 16 43
# 19 Portugal 2015 34.8 1.87 12 43
# 20 Portugal 2016 34.8 1.87 11 43
# 21 Portugal 2017 34.8 1.87 12 43
# 22 Portugal 2018 34.8 1.87 12 43
# 23 Portugal 2019 38.7 1.58 13 43
# 24 Portugal 2020 40 1.5 8.5 43

```

Plot of Portugal's ranking in Europe. Below we reproduce the percentage plot to show how Portugal ranks in relation to six other European countries. Note that the highest percentage is ranked first and the lowest last.

```
cWPx %>%
  filter(Country %in% c("Portugal", "Sweden",
    "Spain", "Hungary", "Romania", "Finland",
    "Germany")) %>%
  ggplot(aes(Year, RankC, colour=Country)) +
  geom_line() +
  geom_point() +
  scale_x_continuous(breaks=seq(1995, 2020, 5)) +
  scale_y_continuous(limits=c(0, 45),
    breaks=seq(0, 45, by=10)) +
  ggtitle("Women in Parliament: Ranked") +
  ylab("Rank in Europe")
```



Interpretation. A total of 38 European countries had data in 1997, 38 in 1997 and 43 in 2020. Within Europe, Portugal was typically ranked in the second quartile (25-50%) with the fluctuations of its ranking most likely due to random variation.

Exercise. How does your chosen country rank within its continent?


Highest rank by year and continent. Which countries have the highest rank in 1997 and 2020? The answer will coincide with the highest percentages (see above).

```
cWPx %>%
  filter(Year %in% c(1997, 2020) & RankC==1) %>%
  arrange(Continent, Year) %>%
  select(Continent, Year, Country, pctWiP, Ratio)
# # A tibble: 10 x 5
# # Groups:   Continent, Year [10]
#   Continent Year Country      pctWiP Ratio
#   <chr>      <dbl> <chr>      <dbl> <dbl>
# 1 Africa    1997 Seychelles  27.3  2.67
# 2 Africa    2020 Rwanda     61.2  0.633
# 3 Americas  1997 Argentina  27.6  2.62
# 4 Americas  2020 Cuba     53.2  0.879
# 5 Asia      1997 Vietnam  26.2  2.81
# 6 Asia      2020 United Arab ~ 50    1
# 7 Europe    1997 Sweden   40.4  1.48
# 8 Europe    2020 Sweden   47.0  1.13
# 9 Oceania   1997 New Zealand 29.2  2.43
# 10 Oceania  2020 New Zealand 40.8  1.45
```

Overall picture. What are the trends globally? There are various regions defined in the World Bank data. We can plot them and highlight the world “average”.

```
cWP %>%
  filter(is.na(Continent)) %>%
  ggplot(aes(Year, pctWiP, group=Country)) +
  geom_line() +
  gghighlight(Country=="World",
    use_direct_label = FALSE,
    use_group_by = FALSE) +
  scale_x_continuous(breaks=seq(1995, 2020, 5)) +
  scale_y_continuous(limits=c(0, 40),
    breaks=seq(0, 40, by=10)) +
  ggtitle("Women in Parliament: Global Trends") +
  ylab("% Women in Parliament")
```

Women in Parliament: Global Trends



Interpretation. The grey lines show that regardless of how we define region the general trends are upwards. The “World” percentage (black line) increased between 1997 and 2020. In 2020, women in parliament represented 25% (i.e. a ratio of 3 men to each woman), which is still less than half the level before gender parity can be claimed.

10. Conclusion

This guide presented an analysis of the percentage of women in parliament as a real-life case study for some of the tidyverse package. Although the format limited what could be presented, we can conclude that the percentage of women in parliament is increasing but that gender parity in parliaments is still far-off.

There is a lot more that can be said and discussed about the limitations, interpretation and potential impact of this data which the World Bank has nicely summarised.⁷ You are strongly encouraged to read their discussion for a more complete understanding.



⁷ <https://databank.worldbank.org/data/reports.aspx?source=2&type=metadata&series=SG.GEN.PARL.ZS>.